



# Setup Kubernetes Cluster using Kubeadm on Ubuntu

*We are using the following Hostnames & IP Assignments:*

## Kubernetes Master Node

— kube-master : 10.10.121.56

•4 Kubernetes Worker Nodes

—kube-node1: 10.10.121.57

— kube-node2 : 10.10.121.71

— kube-node3 : 10.10.121.84

— kube-node4 : 10.10.121.25

## install Docker on Master (10.10.121.56)

---

*references : <https://docs.docker.com/engine/install/ubuntu/>*

---

## Install kubelet,kubeadm,kubectl and containerd.io on Master and worker nodes

At the end of 2021 kubernetes (release v1.22) plans to end support with docker,

---

*as can be seen on the kubernetes website*  
<https://kubernetes.io/blog/2020/12/02/dont-panic-kubernetes-and-docker/>

---

## install containerd and Configure Private Registry for Kubernetes cluster running with containerd :

```
sudo modprobe overlay
sudo modprobe br_netfilter
cat <<EOF | sudo tee /etc/modules-load.d/containerd.conf
```

```
overlay
br_netfilter
EOF
wget
https://github.com/containerd/containerd/releases/download/v1.7.0/containerd-1.7.0-linux-amd64.tar.gz
tar -xvzf containerd-1.7.0-linux-amd64.tar.gz
cd bin/
cp * /usr/local/bin
cd /etc/systemd/system/
```

---

*For containerd service file : <https://raw.githubusercontent.com/containerd/containerd/main/containerd.service>*

---

```
vi containerd.service
```

```
# Copyright The containerd Authors.
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
#     http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
# implied.
# See the License for the specific language governing permissions and
# limitations under the License.

[Unit]
Description=containerd container runtime
Documentation=https://containerd.io
After=network.target local-fs.target

[Service]
#uncomment to enable the experimental sb service (sandboxed) version of
containerd/cri integration
#Environment="ENABLE_CRI_SANDBOXES=sandboxed"
ExecStartPre=-/sbin/modprobe overlay
ExecStart=/usr/local/bin/containerd

Type=notify
Delegate=yes
KillMode=process
Restart=always
```

```

RestartSec=5
# Having non-zero Limit*s causes performance problems due to accounting
overhead
# in the kernel. We recommend using cgroups to do container-local
accounting.
LimitNPROC=infinity
LimitCORE=infinity
LimitNOFILE=infinity
# Comment TasksMax if your systemd version does not supports it.
# Only systemd 226 and above support this version.
TasksMax=infinity
OOMScoreAdjust=-999

[Install]
WantedBy=multi-user.target
1.

```

**systemctl daemon-reload**

```

wget
https://github.com/opencontainers/runc/releases/download/v1.1.6/runc.am
d64
install -m 755 runc.amd64 /usr/local/sbin/runc
wget
https://github.com/containernetworking/plugins/releases/download/v1.2.0
/cni-plugins-linux-amd64-v1.2.0.tgz
mkdir -p /opt/cni/bin
tar Cxzf /opt/cni/bin cni-plugins-linux-amd64-v1.2.0.tgz
mkdir -p /etc/containerd
containerd config default > /etc/containerd/config.toml
systemctl restart containerd

```

**vi /etc/containerd/config.toml**

```

disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2

[cgroup]
  path = ""

[debug]
  address = ""
  format = ""
  gid = 0
  level = ""

```

```
uid = 0

[grpc]
  address = "/run/containerd/containerd.sock"
  gid = 0
  max_recv_message_size = 16777216
  max_send_message_size = 16777216
  tcp_address = ""
  tcp_tls_ca = ""
  tcp_tls_cert = ""
  tcp_tls_key = ""
  uid = 0

[metrics]
  address = ""
  grpc_histogram = false

[plugins]

[plugins."io.containerd.gc.v1.scheduler"]
  deletion_threshold = 0
  mutation_threshold = 100
  pause_threshold = 0.02
  schedule_delay = "0s"
  startup_delay = "100ms"

[plugins."io.containerd.grpc.v1.cri"]
  cdi_spec_dirs = ["/etc/cdi", "/var/run/cdi"]
  device_ownership_from_security_context = false
  disable_apparmor = false
  disable_cgroup = false
  disable_hugetlb_controller = true
  disable_proc_mount = false
  disable_tcp_service = true
  drain_exec_sync_io_timeout = "0s"
  enable_cdi = false
  enable_selinux = false
  enable_tls_streaming = false
  enable_unprivileged_icmp = false
  enable_unprivileged_ports = false
  ignore_image_defined_volumes = false
  image_pull_progress_timeout = "1m0s"
  max_concurrent_downloads = 3
  max_container_log_line_size = 16384
  netns_mounts_under_state_dir = false
  restrict_oom_score_adj = false
  sandbox_image = "registry.k8s.io/pause:3.9"
  selinux_category_range = 1024
  stats_collect_period = 10
  stream_idle_timeout = "4h0m0s"
  stream_server_address = "127.0.0.1"
  stream_server_port = "0"
  systemd_cgroup = false
  tolerate_missing_hugetlb_controller = true
  unset_seccomp_profile = ""

[plugins."io.containerd.grpc.v1.cri".cni]
  bin_dir = "/opt/cni/bin"
```

```

conf_dir = "/etc/cni/net.d"
conf_template = ""
ip_pref = ""
max_conf_num = 1
setup_serially = false

[plugins."io.containerd.grpc.v1.cri".containerd]
  default_runtime_name = "runc"
  disable_snapshot_annotations = true
  discard_unpacked_layers = false
  ignore_blockio_not_enabled_errors = false
  ignore_rdt_not_enabled_errors = false
  no_pivot = false
  snapshotter = "overlayfs"

[plugins."io.containerd.grpc.v1.cri".containerd.default_runtime]
  base_runtime_spec = ""
  cni_conf_dir = ""
  cni_max_conf_num = 0
  container_annotations = []
  pod_annotations = []
  privileged_without_host_devices = false
  privileged_without_host_devices_all_devices_allowed = false
  runtime_engine = ""
  runtime_path = ""
  runtime_root = ""
  runtime_type = ""
  sandbox_mode = ""
  snapshotter = ""

[plugins."io.containerd.grpc.v1.cri".containerd.default_runtime.options
]

[plugins."io.containerd.grpc.v1.cri".containerd.runtimes]

  [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc]
    base_runtime_spec = ""
    cni_conf_dir = ""
    cni_max_conf_num = 0
    container_annotations = []
    pod_annotations = []
    privileged_without_host_devices = false
    privileged_without_host_devices_all_devices_allowed = false
    runtime_engine = ""
    runtime_path = ""
    runtime_root = ""
    runtime_type = "io.containerd.runc.v2"
    sandbox_mode = "podsandbox"
    snapshotter = ""

[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc.options]
  BinaryName = ""
  CriuImagePath = ""
  CriuPath = ""
  CriuWorkPath = ""
  IoGid = 0

```

```

        IoUid = 0
        NoNewKeyring = false
        NoPivotRoot = false
        Root = ""
        ShimCgroup = ""
        SystemdCgroup = true

[plugins."io.containerd.grpc.v1.cri".containerd.untrusted_workload_runtime]
    base_runtime_spec = ""
    cni_conf_dir = ""
    cni_max_conf_num = 0
    container_annotations = []
    pod_annotations = []
    privileged_without_host_devices = false
    privileged_without_host_devices_all_devices_allowed = false
    runtime_engine = ""
    runtime_path = ""
    runtime_root = ""
    runtime_type = ""
    sandbox_mode = ""
    snapshotter = ""

[plugins."io.containerd.grpc.v1.cri".containerd.untrusted_workload_runtime.options]

[plugins."io.containerd.grpc.v1.cri".image_decryption]
    key_model = "node"

[plugins."io.containerd.grpc.v1.cri".registry]
    config_path = ""

[plugins."io.containerd.grpc.v1.cri".registry.auths]

[plugins."io.containerd.grpc.v1.cri".registry.configs]

[plugins."io.containerd.grpc.v1.cri".registry.configs."10.10.121.29:8085".tls]
    insecure_skip_verify = true

[plugins."io.containerd.grpc.v1.cri".registry.configs."10.10.121.29:8085".auth]
    auth = "*****=="

[plugins."io.containerd.grpc.v1.cri".registry.headers]

[plugins."io.containerd.grpc.v1.cri".registry.mirrors]

[plugins."io.containerd.grpc.v1.cri".registry.mirrors."10.10.121.29:8085"]
    endpoint = ["http://10.10.121.29:8085"]

[plugins."io.containerd.grpc.v1.cri".x509_key_pair_streaming]
    tls_cert_file = ""
    tls_key_file = ""

```

```
[plugins."io.containerd.internal.v1.opt"]
  path = "/opt/containerd"

[plugins."io.containerd.internal.v1.restart"]
  interval = "10s"

[plugins."io.containerd.internal.v1.tracing"]
  sampling_ratio = 1.0
  service_name = "containerd"

[plugins."io.containerd.metadata.v1.bolt"]
  content_sharing_policy = "shared"

[plugins."io.containerd.monitor.v1.cgroups"]
  no_prometheus = false

[plugins."io.containerd.nri.v1.nri"]
  disable = true
  disable_connections = false
  plugin_config_path = "/etc/nri/conf.d"
  plugin_path = "/opt/nri/plugins"
  plugin_registration_timeout = "5s"
  plugin_request_timeout = "2s"
  socket_path = "/var/run/nri/nri.sock"

[plugins."io.containerd.runtime.v1.linux"]
  no_shim = false
  runtime = "runc"
  runtime_root = ""
  shim = "containerd-shim"
  shim_debug = false

[plugins."io.containerd.runtime.v2.task"]
  platforms = ["linux/amd64"]
  sched_core = false

[plugins."io.containerd.service.v1.diff-service"]
  default = ["walking"]

[plugins."io.containerd.service.v1.tasks-service"]
  blockio_config_file = ""
  rdt_config_file = ""

[plugins."io.containerd.snapshotter.v1.aufs"]
  root_path = ""

[plugins."io.containerd.snapshotter.v1.btrfs"]
  root_path = ""

[plugins."io.containerd.snapshotter.v1.devmapper"]
  async_remove = false
  base_image_size = ""
  discard_blocks = false
  fs_options = ""
  fs_type = ""
  pool_name = ""
  root_path = ""
```

```

[plugins."io.containerd.snapshotter.v1.native"]
    root_path = ""

[plugins."io.containerd.snapshotter.v1.overlayfs"]
    root_path = ""
    upperdir_label = false

[plugins."io.containerd.snapshotter.v1.zfs"]
    root_path = ""

[plugins."io.containerd.tracing.processor.v1.otlp"]
    endpoint = ""
    insecure = false
    protocol = ""

[plugins."io.containerd.transfer.v1.local"]

[proxy_plugins]

[stream_processors]

    [stream_processors."io.containerd.ocicrypt.decoder.v1.tar"]
        accepts = ["application/vnd.oci.image.layer.v1.tar+encrypted"]
        args = ["--decryption-keys-path", "/etc/containerd/ocicrypt/keys"]
        env =
["OCICRYPT_KEYPROVIDER_CONFIG=/etc/containerd/ocicrypt/ocicrypt_keyprov
ider.conf"]
        path = "ctd-decoder"
        returns = "application/vnd.oci.image.layer.v1.tar"

    [stream_processors."io.containerd.ocicrypt.decoder.v1.tar.gzip"]
        accepts = ["application/vnd.oci.image.layer.v1.tar+gzip+encrypted"]
        args = ["--decryption-keys-path", "/etc/containerd/ocicrypt/keys"]
        env =
["OCICRYPT_KEYPROVIDER_CONFIG=/etc/containerd/ocicrypt/ocicrypt_keyprov
ider.conf"]
        path = "ctd-decoder"
        returns = "application/vnd.oci.image.layer.v1.tar+gzip"

[timeouts]
    "io.containerd.timeout.bolt.open" = "0s"
    "io.containerd.timeout.metrics.shimstats" = "2s"
    "io.containerd.timeout.shim.cleanup" = "5s"
    "io.containerd.timeout.shim.load" = "5s"
    "io.containerd.timeout.shim.shutdown" = "3s"
    "io.containerd.timeout.task.state" = "2s"

[ttrpc]
    address = ""
    gid = 0
    uid = 0

```

**Note :** you need to replace ip and token



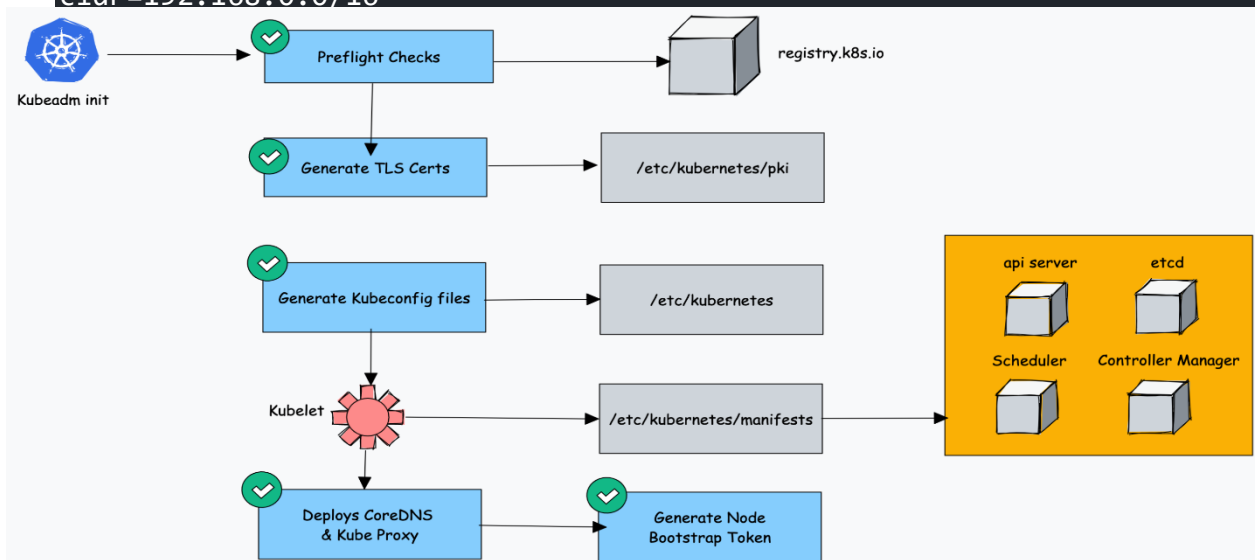
## Install Kubernetes

```
swapoff -a
```

Kubernetes references : <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

### In Master node

```
kubeadm init --apiserver-advertise-address=10.10.121.56 --pod-network-cidr=192.168.0.0/16
```



### Message after Kubeadm init is successful:

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "`kubectl apply -f [podnetwork].yaml`" with one of the options listed at:

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 10.10.121.56:6443 --token zxen48.z5ruut0pwf8rxb4i --
```

To View the Join command:

```
kubeadm token create --print-join-command
```

Join the workers with join command.

Install flannel Network Plugin for Pod Networking:

Reference : <https://kubernetes.io/docs/concepts/cluster-administration/addons/#networking-and-network-policy>

```
wget https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
```

```
kubectl apply -f kube-flannel
```

Note 1 : Before above 'apply' command, we have to check CIDR range of cluster in kube-flannel.yaml file. If it does not match with that of mentioned in 'kubeadm init' command.....we have to change it.

Node 2 : We have to wait till all the nodes get into 'Ready' state.

Output:

```
root@kube-master:~# kubectl get nodes
NAME           STATUS    ROLES    AGE   VERSION
kube-master    Ready     control-plane   9d    v1.27.1
kube-node1     Ready     <none>        9d    v1.27.1
kube-node2     Ready     <none>        9d    v1.27.1
kube-node3     Ready     <none>        9d    v1.27.1
kube-node4     Ready     <none>        9d    v1.27.1
```

Install metal lb :

Reference : <https://metallb.universe.tf/installation/>

```
kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.13.9/config/manifests/metallb-native.yaml
```

```
vi l2advertisement.yaml
```

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: example
  namespace: metallb-system
spec:
  ipAddressPools:
```

```
- first-pool
```

```
vi ipaddresspoll.yaml
```

```
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  name: first-pool
  namespace: metallb-system
spec:
  addresses:
    - 10.10.XXXX-10.10.XXXXX
```

```
kubectl create -f ipaddresspoll.yaml
```

```
kubectl create -f l2advertisement.yaml
```

**Install Helm:**

```
wget https://get.helm.sh/helm-v3.11.2-linux-amd64.tar.gz
```

```
tar -zxvf helm-v3.11.2-linux-amd64.tar.gz
```

```
mv linux-amd64/helm /usr/local/bin/helm
```

**Install Ingress :**

```
git clone https://github.com/nginxinc/kubernetes-ingress.git --branch v3.1.0
```

```
cd kubernetes-ingress/deployments/helm-chart
```

```
helm repo add nginx-stable https://helm.nginx.com/stable
kubectl apply -f crds/
```

```
helm install nginx-controller-one nginx-stable/nginx-ingress --set
controller.ingressClass=nginx-one --namespace ingress-nginx-one --
create-namespace
```

```
kubectl get svc -n ingress-nginx-one
```